

# Evaluation Metrics Overview

Accuracy measures the proportion of correct predictions across all classes.

Precision focuses on the quality of positive predictions, limiting false alarms.

Recall captures coverage of actual positives, ensuring critical cases are caught.

The F1 score balances precision and recall, rewarding joint improvements.

ROC-AUC evaluates ranking quality over thresholds, illustrating trade-offs.

PR-AUC highlights positive class performance under imbalance.

Calibration curves reveal whether predicted probabilities align with outcomes.